

MOVIE RECOMMENDATION SYSTEM USING ML

In [1]:

```
#IMPORTING THE REQUIRED LIBRARIES -
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O
import warnings
warnings.filterwarnings("ignore")

#VISUALISATION
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

#MACHINE LEARNING MODEL
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
```

In [2]:

```
#LOADING THE REQUIRED DATASETS
df_movies= pd.read_csv('movies.csv') #contains data about movies
df_rating=pd.read_csv('ratings.csv') #contains data about ratings

#Since user data doesn't have any column name we are allotting column name.
column_names_users = ['user_id', 'gender', 'age', 'occupation', 'zipcode']
df_user= pd.read_csv('users.csv', delimiter = ';', names= column_names_users,header = 0)
df_user.dropna(inplace=True) #dropping na values
```

In [3]:

```
#dropping unwanted columns from movies dataset
df_movie=df_movies.drop('genres',axis=1)
```

In [4]:

```
#merging movies and ratings dataset on column movie id
movie_rating_df=pd.merge(df_movie,df_rating,on="movieId")
movie_rating_df.head()
```

Out[4]:

| | movieId | title | userId | rating | timestamp |
|---|---------|------------------|--------|--------|------------|
| 0 | 1 | Toy Story (1995) | 1 | 4.0 | 964982703 |
| 1 | 1 | Toy Story (1995) | 5 | 4.0 | 847434962 |
| 2 | 1 | Toy Story (1995) | 7 | 4.5 | 1106635946 |
| 3 | 1 | Toy Story (1995) | 15 | 2.5 | 1510577970 |
| 4 | 1 | Toy Story (1995) | 17 | 4.5 | 1305696483 |

In [5]:

```
#creating a table with all needed columns
df = pd.concat([movie_rating_df,df_user], axis=1)
df.head()
```

Out[5]:

| | movieId | title | userId | rating | timestamp | user_id | gender | age | occupation | zipcode |
|---|---------|------------------|--------|--------|------------|---------|--------|------|------------|---------|
| 0 | 1 | Toy Story (1995) | 1 | 4.0 | 964982703 | 1.0 | F | 1.0 | 10.0 | 48067 |
| 1 | 1 | Toy Story (1995) | 5 | 4.0 | 847434962 | 2.0 | M | 56.0 | 16.0 | 70072 |
| 2 | 1 | Toy Story (1995) | 7 | 4.5 | 1106635946 | 3.0 | M | 25.0 | 15.0 | 55117 |
| 3 | 1 | Toy Story (1995) | 15 | 2.5 | 1510577970 | 4.0 | M | 45.0 | 7.0 | 2460 |
| 4 | 1 | Toy Story (1995) | 17 | 4.5 | 1305696483 | 5.0 | M | 25.0 | 20.0 | 55455 |

In [6]:

```
df.isna().any() #checking for null values
```

Out[6]:

```
movieId      False
title        False
userId       False
rating       False
timestamp    False
user_id      True
gender       True
age          True
occupation   True
zipcode      True
dtype: bool
```

In [7]:

```
df=df.dropna() #dropping the null values
```

In [8]:

```
#ADDING NEW COLUMN
```

```
#to create seperate column for age group data
```

```
labels = ['0-9', '10-19', '20-29', '30-39', '40-49', '50-59', '60-69', '70-79']
df['age_group'] = pd.cut(df.age, range(0, 81, 10), right=False, labels=labels)
df[['age', 'age_group']].drop_duplicates()[:10]
```

Out[8]:

| | age | age_group |
|----|------|-----------|
| 0 | 1.0 | 0-9 |
| 1 | 56.0 | 50-59 |
| 2 | 25.0 | 20-29 |
| 3 | 45.0 | 40-49 |
| 5 | 50.0 | 50-59 |
| 6 | 35.0 | 30-39 |
| 17 | 18.0 | 10-19 |

In [9]:

```
# We need to categorize the imdb values in the range of 1,2,3,4 and 5 to mark them as the bad,average,good and excellent movies r
df["rating"]=pd.cut(df['rating'], bins=[1,2,3,4,5], right=True, labels=False)+1
```

In [10]:

```
#First 500 extracted records
df.dropna(inplace=True)
```

In [11]:

```
df['gender']=df['gender'].replace('M',1) #CONVERTING CATEGORICAL VALUE TO BINARY VALUE
df['gender']=df['gender'].replace('F',0)
```

EDA

In [12]:

```
df.describe()
```

Out[12]:

| | movieId | userId | rating | timestamp | user_id | gender | age | occupation |
|-------|-------------|-------------|-------------|--------------|-------------|-------------|-------------|-------------|
| count | 5748.000000 | 5748.000000 | 5748.000000 | 5.748000e+03 | 5748.000000 | 5748.000000 | 5748.000000 | 5748.000000 |
| mean | 92.287752 | 310.526444 | 2.692241 | 1.052242e+09 | 2997.955115 | 0.716249 | 30.740953 | 8.164057 |
| std | 69.166061 | 185.329731 | 0.918705 | 2.200185e+08 | 1740.336688 | 0.450857 | 12.860763 | 6.332890 |
| min | 1.000000 | 1.000000 | 1.000000 | 8.281246e+08 | 1.000000 | 0.000000 | 1.000000 | 0.000000 |
| 25% | 31.000000 | 142.000000 | 2.000000 | 8.474352e+08 | 1494.750000 | 0.000000 | 25.000000 | 3.000000 |
| 50% | 81.000000 | 313.000000 | 3.000000 | 9.731055e+08 | 2996.500000 | 1.000000 | 25.000000 | 7.000000 |
| 75% | 158.000000 | 469.000000 | 3.000000 | 1.189457e+09 | 4491.250000 | 1.000000 | 35.000000 | 14.000000 |
| max | 223.000000 | 610.000000 | 4.000000 | 1.537799e+09 | 6040.000000 | 1.000000 | 56.000000 | 20.000000 |

In [13]:

```
df.shape
```

Out[13]:

(5748, 11)


In [14]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5748 entries, 0 to 6039
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   movieId     5748 non-null   int64
1   title       5748 non-null   object
2   userId      5748 non-null   int64
3   rating      5748 non-null   float64
4   timestamp   5748 non-null   int64
5   user_id     5748 non-null   float64
6   gender      5748 non-null   int64
7   age         5748 non-null   float64
8   occupation  5748 non-null   float64
9   zipcode     5748 non-null   object
10  age_group   5748 non-null   category
dtypes: category(1), float64(4), int64(4), object(2)
memory usage: 499.9+ KB
```

In [15]:

```
#Visualize user age distribution
df['age'].value_counts().plot(kind='barh',alpha=0.7,figsize=(10,10))
plt.show()
```



In [16]:

```
df
```

Out[16]:

| | movieId | title | userId | rating | timestamp | user_id | gender | age | occupation | zipcode | age_group |
|------|---------|------------------|--------|--------|------------|---------|--------|------|------------|---------|-----------|
| 0 | 1 | Toy Story (1995) | 1 | 3.0 | 964982703 | 1.0 | 0 | 1.0 | 10.0 | 48067 | 0-9 |
| 1 | 1 | Toy Story (1995) | 5 | 3.0 | 847434962 | 2.0 | 1 | 56.0 | 16.0 | 70072 | 50-59 |
| 2 | 1 | Toy Story (1995) | 7 | 4.0 | 1106635946 | 3.0 | 1 | 25.0 | 15.0 | 55117 | 20-29 |
| 3 | 1 | Toy Story (1995) | 15 | 2.0 | 1510577970 | 4.0 | 1 | 45.0 | 7.0 | 2460 | 40-49 |
| 4 | 1 | Toy Story (1995) | 17 | 4.0 | 1305696483 | 5.0 | 1 | 25.0 | 20.0 | 55455 | 20-29 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6035 | 223 | Clerks (1994) | 591 | 3.0 | 970524802 | 6036.0 | 0 | 25.0 | 15.0 | 32603 | 20-29 |
| 6036 | 223 | Clerks (1994) | 594 | 3.0 | 1109037031 | 6037.0 | 0 | 45.0 | 1.0 | 76006 | 40-49 |
| 6037 | 223 | Clerks (1994) | 597 | 2.0 | 941558145 | 6038.0 | 0 | 56.0 | 1.0 | 14706 | 50-59 |
| 6038 | 223 | Clerks (1994) | 599 | 3.0 | 1498500829 | 6039.0 | 0 | 45.0 | 0.0 | 1060 | 40-49 |
| 6039 | 223 | Clerks (1994) | 600 | 2.0 | 1237715120 | 6040.0 | 1 | 25.0 | 6.0 | 11106 | 20-29 |

5748 rows × 11 columns

SPLITTING THE DATASET FOR MACHINE LEARNING MODEL

```
#Use the following features:movie id,age,occupation
x = df.drop('rating',axis=1) #x=independent
x.drop(['title','userId','occupation','zipcode','age_group'],axis=1,inplace=True)
```

In [72]:

```
#features:gender,age,occupation
x = df.drop('rating',axis=1) #x=independent
x.drop(['user_id','zipcode','age_group','timestamp','movieId','userId','title'],axis=1,inplace=True)
```

In [73]:

x

Out[73]:

| | gender | age | occupation |
|------|--------|------|------------|
| 0 | 0 | 1.0 | 10.0 |
| 1 | 1 | 56.0 | 16.0 |
| 2 | 1 | 25.0 | 15.0 |
| 3 | 1 | 45.0 | 7.0 |
| 4 | 1 | 25.0 | 20.0 |
| ... | ... | ... | ... |
| 6035 | 0 | 25.0 | 15.0 |
| 6036 | 0 | 45.0 | 1.0 |
| 6037 | 0 | 56.0 | 1.0 |
| 6038 | 0 | 45.0 | 0.0 |
| 6039 | 1 | 25.0 | 6.0 |

5748 rows × 3 columns

In [78]:

```
#Used rating as Label
y =df['rating'].values #dependent= rating
df['rating']
```

Out[78]:

```
0      3.0
1      3.0
2      4.0
3      2.0
4      4.0
...
6035   3.0
6036   3.0
6037   2.0
6038   3.0
6039   2.0
Name: rating, Length: 5748, dtype: float64
```

In [77]:

y

Out[77]:

array([3., 3., 4., ..., 2., 3., 2.])

In [75]:

```
#Create train and test data set
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=42)
```

In [76]:

```
# Decision Tree

decision_tree = DecisionTreeClassifier()
decision_tree.fit(x_train,y_train)
Y_pred = decision_tree.predict(x_test)
Y_pred
```

Out[76]:

array([3., 3., 3., ..., 3., 3., 3.])

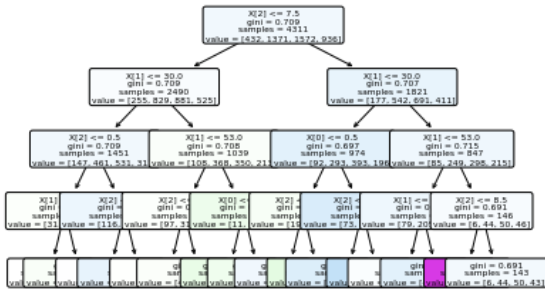
In [61]:

```
import sklearn.datasets as datasets
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

import sklearn
print(sklearn.__version__)

from sklearn import tree
clf=DecisionTreeClassifier(max_depth=4,random_state=0)
clf.fit(x_train,y_train)
tree.plot_tree(clf,fontsize=6,filled=True,
               rounded=True)
plt.figure(figsize=(10,10))
plt.savefig('mov.png')
from matplotlib import pyplot as plt
```

1.0.2



<Figure size 720x720 with 0 Axes>

In [62]:

```
acc_dec = round(decision_tree.score(x_train,y_train) * 100, 2)
acc_dec
```

Out[62]:

41.68

GAUSSIAN NAIVE BAYES

In [63]:

```
# Gaussian Naive Bayes

gaussian = GaussianNB()
gauss=gaussian.fit(x_train,y_train)
Y_pred = gaussian.predict(x_test)
Y_pred
```

Out[63]:


array([2., 3., 3., ..., 3., 3., 3.])

In [64]:

```
from sklearn.metrics import plot_confusion_matrix
plot_confusion_matrix(gauss,x_test,y_test)
```

Out[64]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2a4c93fcd30>
```



In [65]:

```
acc_gaussian = round(gaussian.score(x_train,y_train) * 100, 2)
acc_gaussian
```

Out[65]:

36.63

KNN

In [69]:

```
# K Nearest Neighbors Classifier
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(x_train, y_train)
Y_predic = knn.predict(x_test)
Y_predic
```

Out[69]:

```
array([1., 3., 4., ..., 2., 3., 2.])
```

In [33]:

#KNN PLOTTING

In [70]:

```
acc_knn = round(knn.score(x_train,y_train) * 100, 2)
acc_knn
```

Out[70]:

33.66

In [71]:

```
models = pd.DataFrame({
    'Model': [ 'KNN', 'Random Forest', 'Naive Bayes', 'Decision Tree'],
    'Score': [ acc_knn, acc_random_forest, acc_gaussian, acc_dec]})
models.sort_values(by='Score', ascending=False)
```

Out[71]:

| | Model | Score |
|---|---------------|-------|
| 1 | Random Forest | 41.68 |
| 3 | Decision Tree | 41.68 |
| 2 | Naive Bayes | 36.63 |
| 0 | KNN | 33.66 |

In [79]:

```
df.groupby('title')['rating'].count().sort_values(ascending=False).head()
```

Out[79]:

```
title
Braveheart (1995)      232
Toy Story (1995)       214
Usual Suspects, The (1995) 202
Apollo 13 (1995)       200
Seven (a.k.a. Se7en) (1995) 199
Name: rating, dtype: int64
```

In [39]:

```
ratings = pd.DataFrame(df.groupby('title')['rating'].mean())
ratings.head()
```

Out[39]:

| | rating |
|---------------------------------------|----------|
| title | |
| Ace Ventura: When Nature Calls (1995) | 2.076923 |
| Addiction, The (1995) | 1.000000 |
| Amateur (1994) | 1.000000 |
| Amazing Panda Adventure, The (1995) | 2.333333 |
| American President, The (1995) | 2.782609 |

In [40]:

```
ratings['num of ratings'] = pd.DataFrame(df.groupby('title')['rating'].count())
ratings.head()
```

Out[40]:

| | rating | num of ratings |
|---------------------------------------|----------|----------------|
| title | | |
| Ace Ventura: When Nature Calls (1995) | 2.076923 | 78 |
| Addiction, The (1995) | 1.000000 | 1 |
| Amateur (1994) | 1.000000 | 1 |
| Amazing Panda Adventure, The (1995) | 2.333333 | 3 |
| American President, The (1995) | 2.782609 | 69 |

In [41]:

```
corr_.sort_values('Correlation',ascending=False).head(n)
```

Out[41]:

| Correlation | |
|--|-----|
| title | |
| Rare Exports: A Christmas Tale (Rare Exports) (2010) | 1.0 |
| Angel-A (2005) | 1.0 |
| Poison Ivy (1992) | 1.0 |
| Dr. Jekyll and Mr. Hyde (1931) | 1.0 |
| Police Academy 6: City Under Siege (1989) | 1.0 |

In [42]:

```
corr_ = corr_.join(ratings['num of ratings'])
corr_.head()
```

Out[42]:

| Correlation | | num of ratings |
|--------------------------------|-----------|----------------|
| title | | |
| 'burbs, The (1989) | 0.034329 | NaN |
| (500) Days of Summer (2009) | 0.352152 | NaN |
| *batteries not included (1987) | 0.419314 | NaN |
| ...And Justice for All (1979) | -1.000000 | NaN |
| 10 Cloverfield Lane (2016) | 0.188982 | NaN |

In [45]:

```
rec=corr_[corr_['num of ratings']>100].sort_values('Correlation',ascending=False).head()
```

In [47]:

```
rec['num of ratings']
```

Out[47]:

| | |
|--------------------------------------|-------|
| title | |
| Apollo 13 (1995) | 200.0 |
| Net, The (1995) | 102.0 |
| Batman Forever (1995) | 128.0 |
| Crimson Tide (1995) | 101.0 |
| Babe (1995) | 124.0 |
| Name: num of ratings, dtype: float64 | |